

Operating Systems

Lecture 03

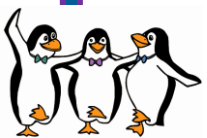
Processes

Dr. Khalid A. Hafeez



Processes

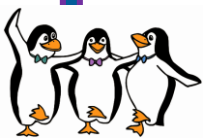
- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems





Objectives

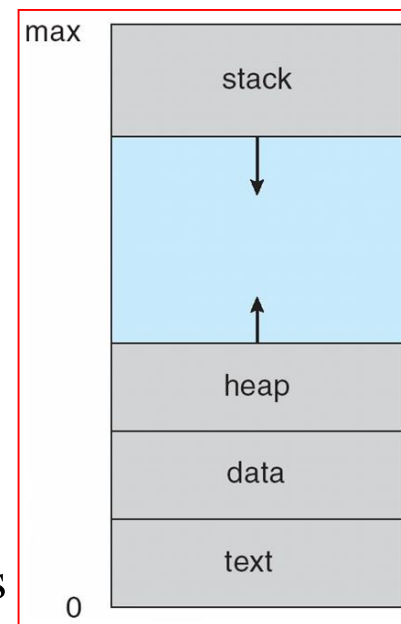
- To introduce the notion of a **process**: a program in execution, which forms the basis of all computation
- To describe the various features of processes, including **scheduling**, **creation** and **termination**, and **communication**
- To explore interprocess communication using shared memory and message passing
- To describe communication in client-server systems



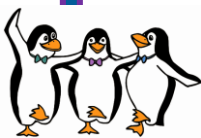


Process Concept

- An operating system executes a variety of programs, what to call all the CPU activities:
 - **Batch system**: executes **jobs**
 - **Time-shared systems**: executes **user programs** or **tasks**
- *Textbook uses the terms **job** and **process** almost interchangeably*
- **Process**: it is a program in execution;
 - process execution must progress in sequential fashion
- **The process has multiple parts other than the code:**
 - The program code, also called **text section**
 - Current activity including **program counter**, processor **registers**
 - **Stack** containing temporary data
 - Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time



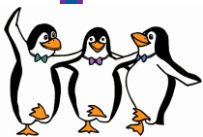
Process in Memory





Process Concept

- Program is a **passive** entity stored on disk (**executable file**), but the process is an **active** entity with program counter pointing to the next instruction to be executed
 - Program becomes process when executable file loaded into memory
- Execution of program can be started via:
 - GUI mouse clicks, or
 - command line entry of its name, etc
- One program can be of several processes
 - Consider multiple users executing the same program
 - Or a user invoking many copies of the web browser

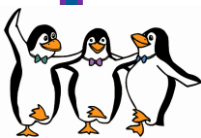
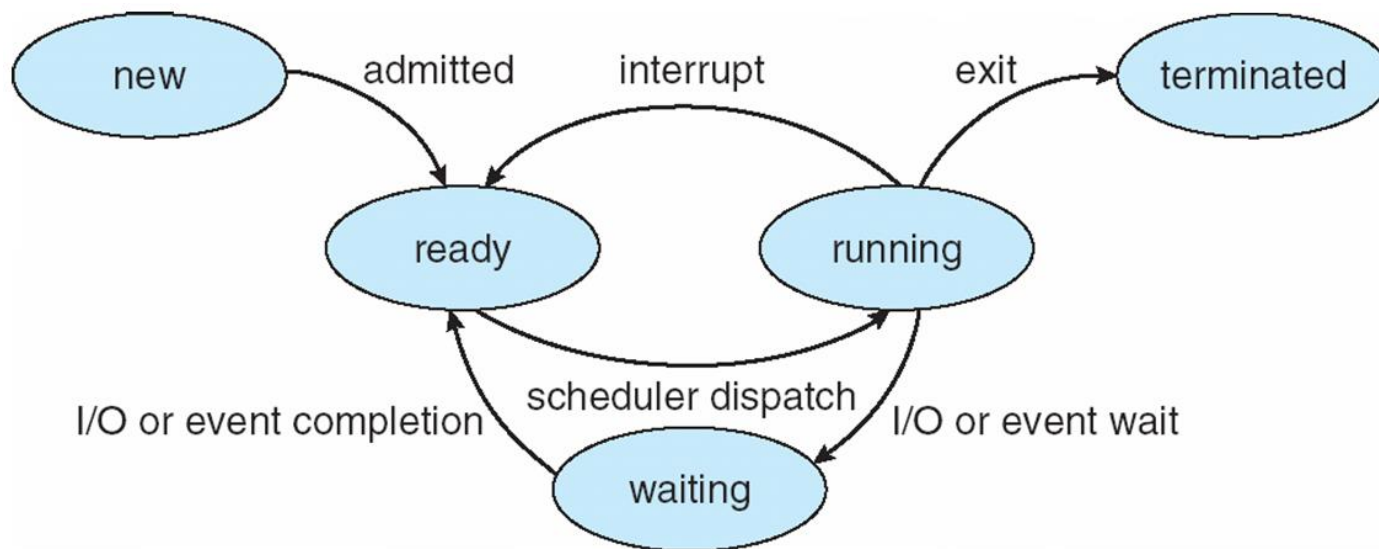




Process Concept

■ Process State

- As a process executes, it changes **state**
- A process may be in one of the following states:
 - ▶ **New**: The process is being created
 - ▶ **Running**: Instructions are being executed
 - ▶ **Waiting**: The process is waiting for some event to occur (such as an I/O completion or reception of a signal).
 - ▶ **Ready**: The process is waiting to be assigned to a processor
 - ▶ **Terminated**: The process has finished execution



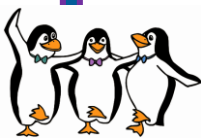


Process Concept

■ Process Control Block (PCB)

- Each process is represented in the OS by a **process control block (PCB)**, also called **task control block**.
- PCB contains the following information:
 - ▶ **Process state**: ready, running, waiting, etc
 - ▶ **Program counter**: location of next instruction to be executed
 - ▶ **CPU registers**: contents of all process-centric registers
 - ▶ **CPU scheduling information**: priorities, scheduling queue pointers
 - ▶ **Memory-management information**: memory allocated to the process
 - ▶ **Accounting information**: amount of CPU used, clock time elapsed since start, time limits, process numbers
 - ▶ **I/O status information**: list of I/O devices allocated to process, list of open files

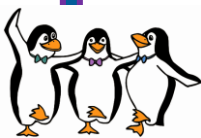
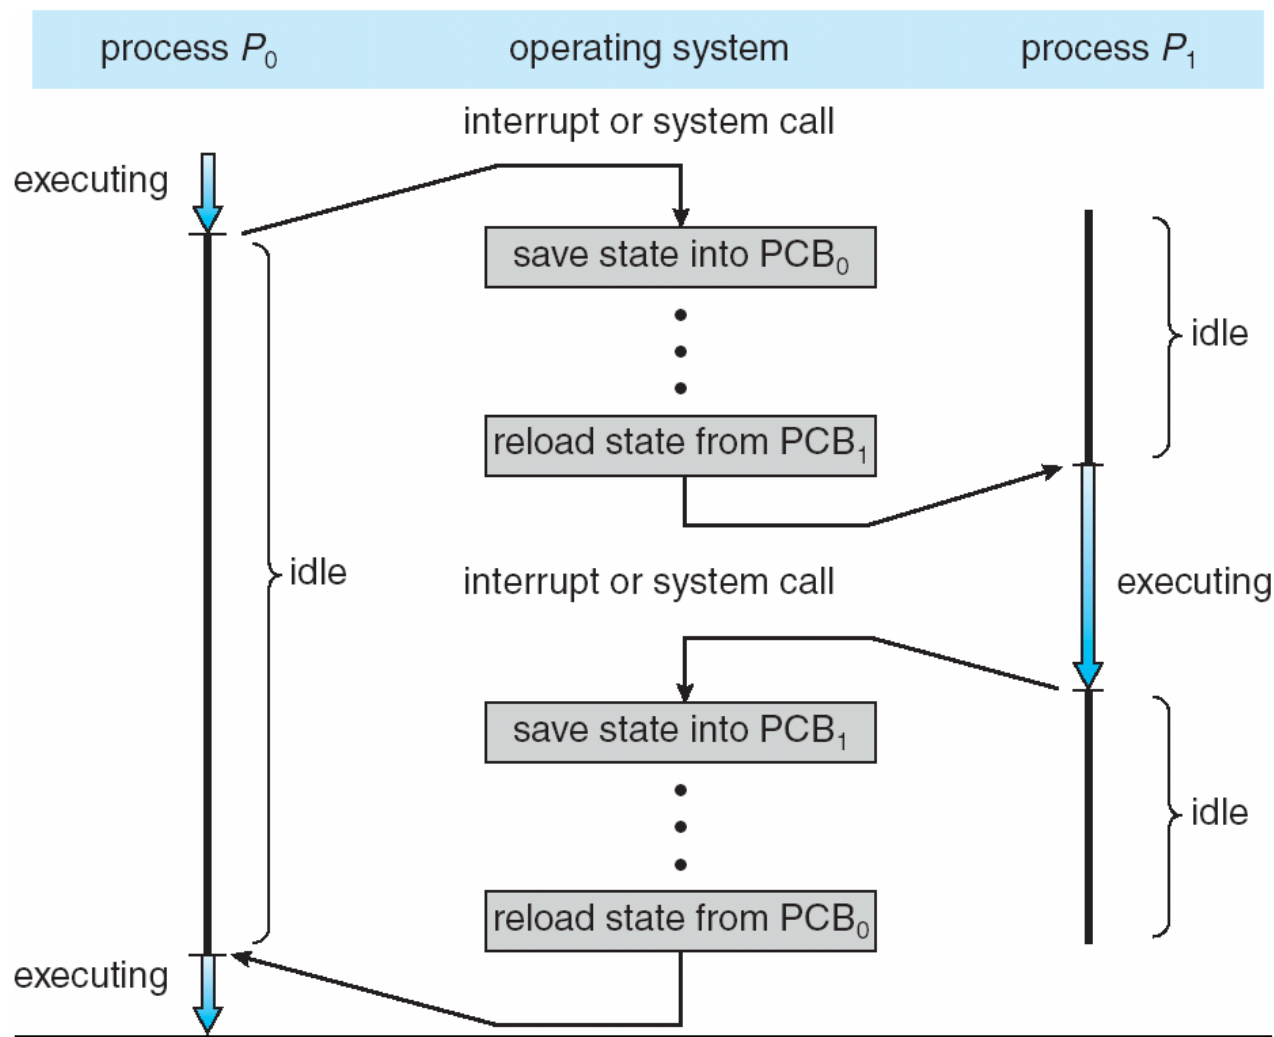
process state
process number
program counter
registers
memory limits
list of open files
...





Process Concept

- CPU Switch From Process to Process

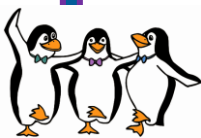




Process Concept

■ Threads

- So far, a process is a program that performs a single **thread** of execution.
- Most OSes allow a process to have multiple threads of execution and thus to perform more than one task at a time.
 - ▶ It is beneficial on multicore systems, where multiple threads can run in parallel.
 - ▶ Consider having multiple program counters per process
 - Then multiple locations can execute at once
 - » Multiple threads of control -> **threads**
 - ▶ The PCB is expanded to include information for each thread, such as multiple program counters



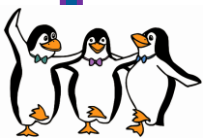
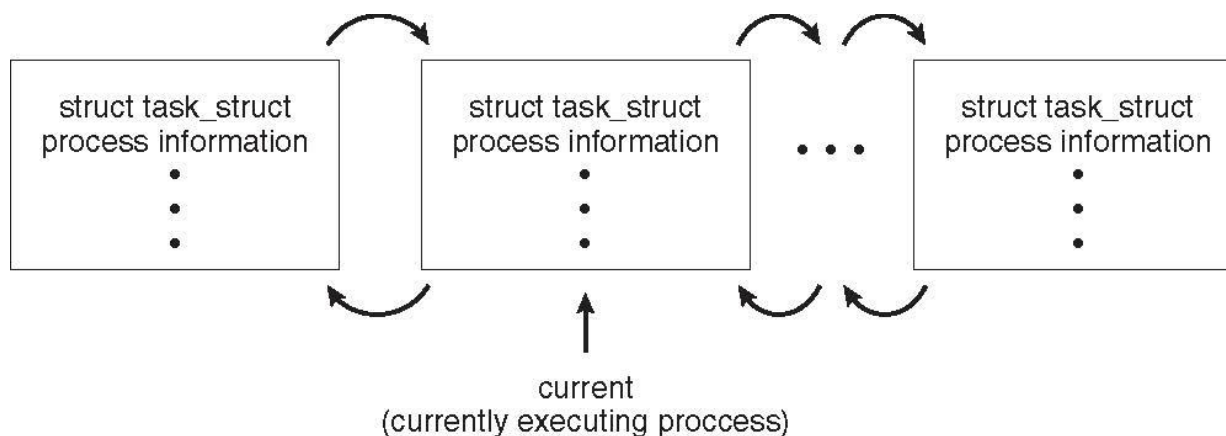


Process Concept

■ Process Representation in Linux

Represented by the C structure `task_struct`

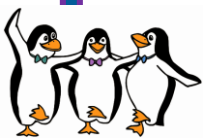
```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```





Process Scheduling

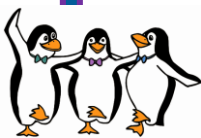
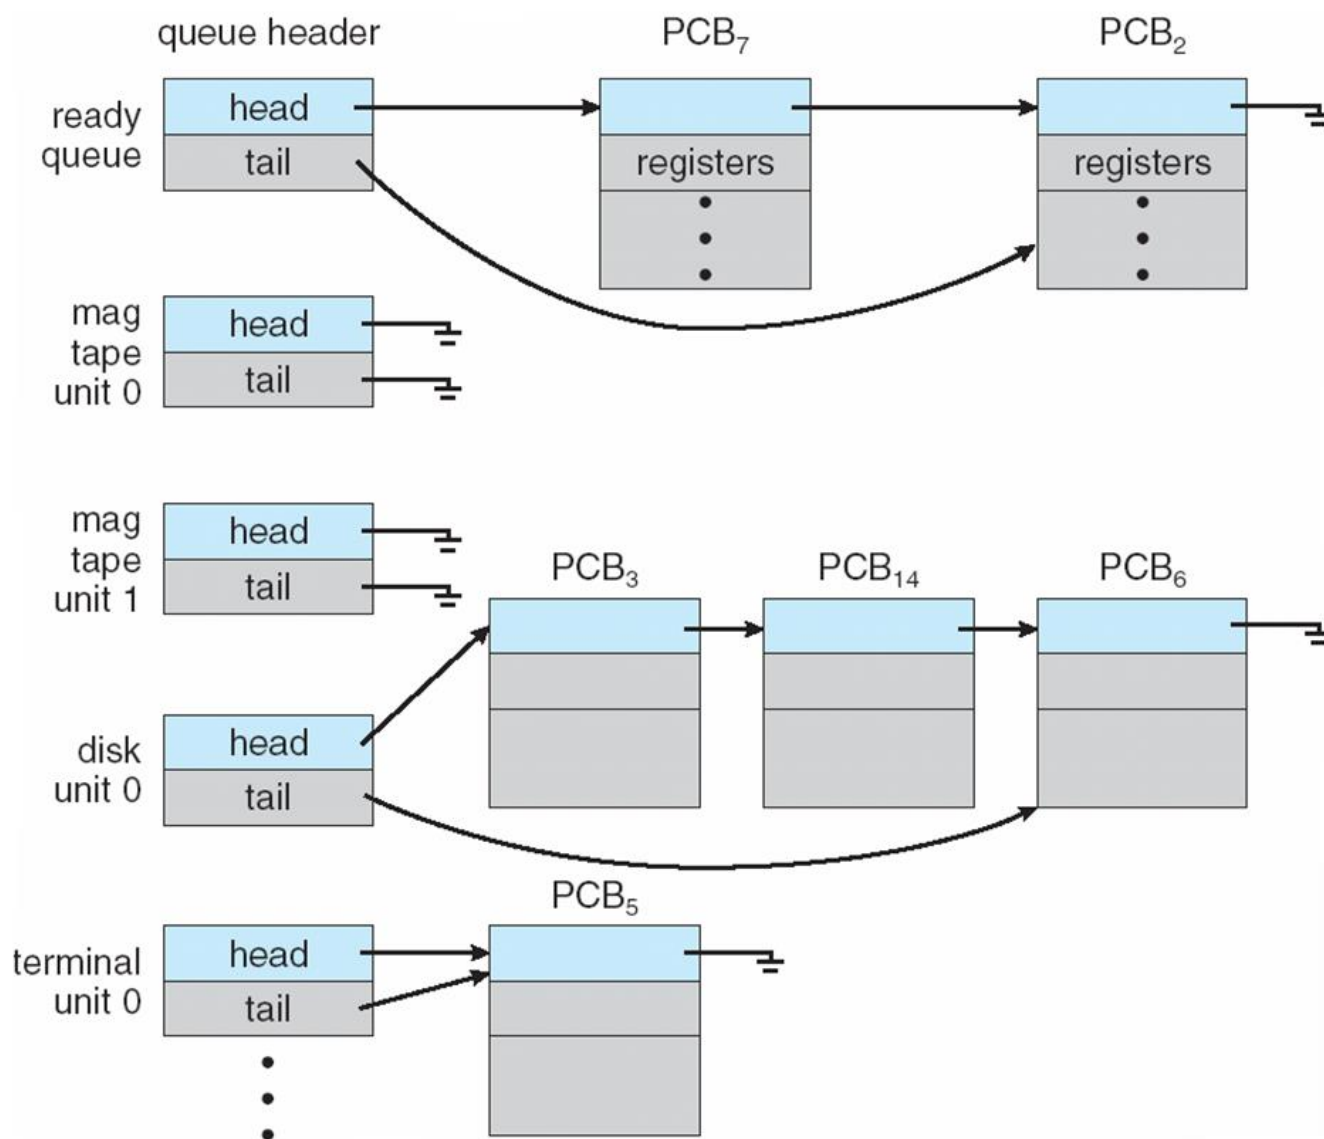
- The objective of **multiprogramming** is to have some process running at all times, to maximize CPU utilization.
- The objective of **time sharing** is to switch the CPU among processes
- For a single-processor system, there will never be more than one running process.
- **Process scheduler** selects among available processes for next execution on CPU
 - Maintains **scheduling queues** of processes
 - ▶ **Job queue**; set of all processes in the system
 - ▶ **Ready queue**: set of all processes residing in main memory, ready and waiting to execute
 - ▶ **Device queues**: set of processes waiting for an I/O device
 - Each device has its own device queue
 - Processes migrate among the various queues





Process Scheduling

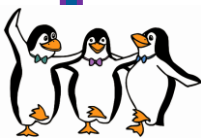
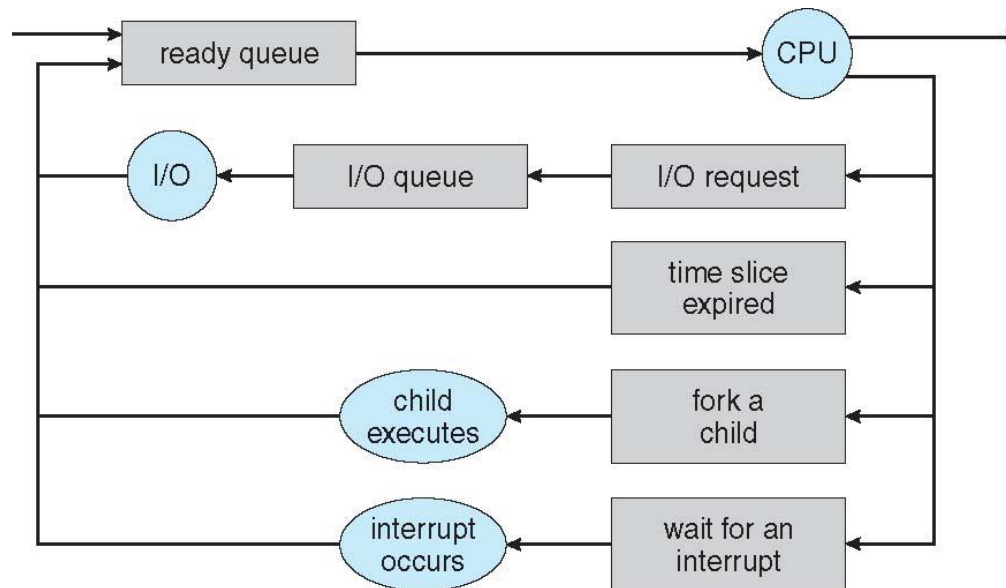
- Ready Queue And Various I/O Device Queues





Process Scheduling

- **Queueing diagram:** is a common representation of process scheduling
 - Each rectangular box represents a queue
 - The circles represent the resources that serve the queues,
 - The arrows indicate the flow of processes in the system.
- A new process is initially put in the ready queue. It waits there until it is selected for execution, one of several events could occur while executing:
 - The process issues an I/O request, then placed in an I/O queue.
 - The process creates a new child process, wait for the child's termination.
 - The process is removed from the CPU, as a result of an interrupt.

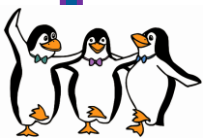




Process Scheduling

■ Schedulers

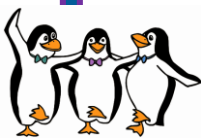
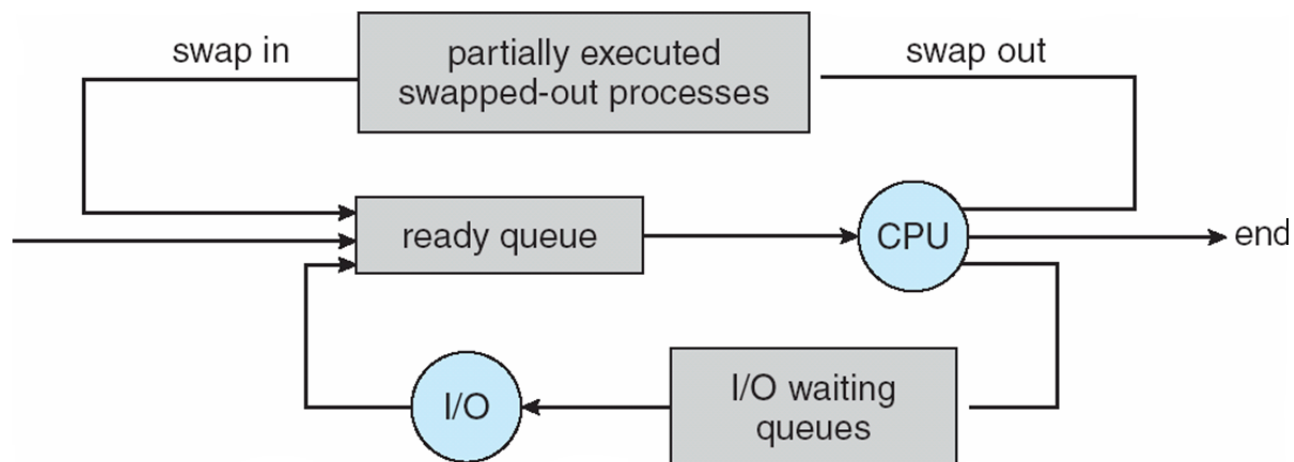
- **Short-term scheduler** (or **CPU scheduler**) – selects which process from ready processes should be executed next and allocates CPU
 - ▶ Sometimes the only scheduler in a system
 - ▶ Short-term scheduler is invoked frequently (milliseconds) \Rightarrow (must be fast)
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
 - ▶ Long-term scheduler is invoked infrequently (seconds, minutes) \Rightarrow (may be slow)
 - ▶ The long-term scheduler controls the **degree of multiprogramming** (the number of processes in memory).
- **Processes can be described as either:**
 - ▶ **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - ▶ **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- The long-term scheduler should select a good **process mix** of I/O-bound and CPU-bound processes.





Process Scheduling

- Some operating systems, such as time-sharing systems, may introduce an additional, intermediate level of scheduling
 - **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store it on disk, bring it back in from disk to continue execution: this is called **swapping**
 - Swapping may be necessary to improve the process mix

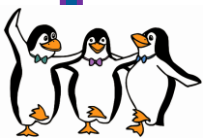




Process Scheduling

■ Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - ▶ The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support
 - ▶ Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once

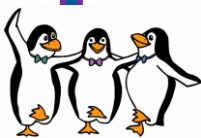




Process Scheduling

■ Multitasking in Mobile Systems

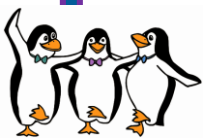
- Some mobile systems (e.g., early version of iOS) allow only one process to run, others are suspended
- Apple now provides a limited form of multitasking for user applications:
 - ▶ Single **foreground** process- controlled via user interface
 - ▶ Multiple **background** processes– in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - ▶ Background process uses a **service** to perform tasks
 - ▶ Service can keep running even if background process is suspended
 - ▶ Service has no user interface, small memory use





Operations on Processes

- System must provide mechanisms for:
 - process creation,
 - process termination,
- The processes in most systems can execute concurrently, and they may be created and deleted dynamically

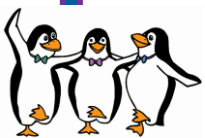
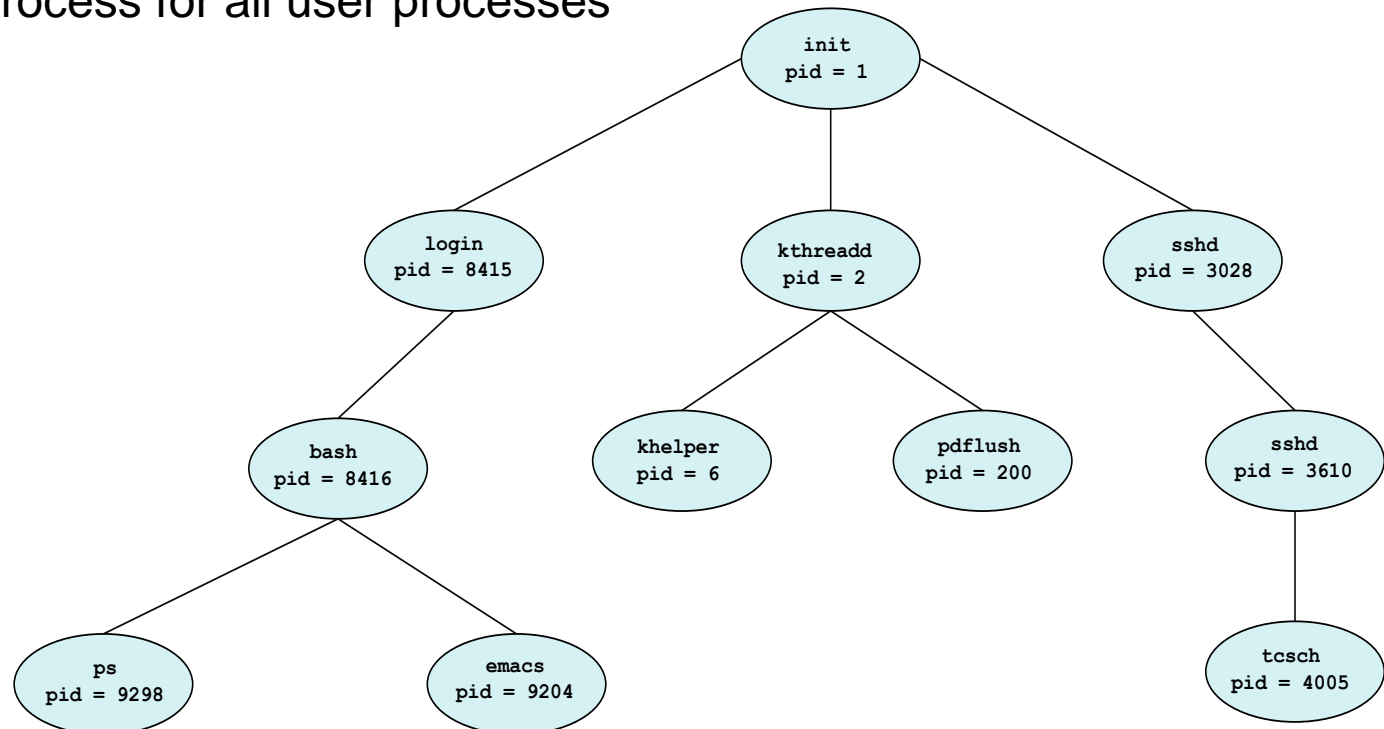




Operations on Processes

■ Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- A Tree of Processes in Linux
 - ▶ The ***init*** process (which always has a pid of 1) serves as the root parent process for all user processes





Operations on Processes

■ Process Creation

- Resource sharing options

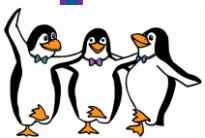
- ▶ A child process may obtain its resources directly from the OS,
- ▶ A child process may share subset of parent's resources
 - This prevents any process from overloading the system by creating too many child processes

- When a process creates a new process, two possibilities for execution exist:

1. The parent continues to execute concurrently with its children.
2. The parent waits until some or all of its children have terminated.

- There are also two address-space possibilities for the new process:

1. The child process is a duplicate of the parent process (it has the same program and data as the parent).
2. The child process has a new program loaded into it.



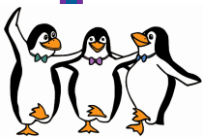


Operations on Processes

■ Process Creation

● UNIX examples

- ▶ *fork()*: system call is for creating a new process
 - The new process consists of a copy of the address space of the original process.
 - Both processes (the parent and the child) continue execution at the instruction after the *fork()*, with one difference:
 - » the return code for the *fork()* is zero for the new (child) process, whereas the (nonzero) process identifier of the child is returned to the parent.
- ▶ *exec()*: system call which is used after a *fork()* to replace the process' memory space with a new program
 - So the two processes are able to communicate and then go their separate ways.
 - The parent can then create more children; or, if it may issue a *wait()* system call to move itself off the ready queue until the termination of the child





Operations on Processes

- C Program Forking Separate Process

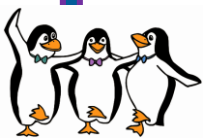
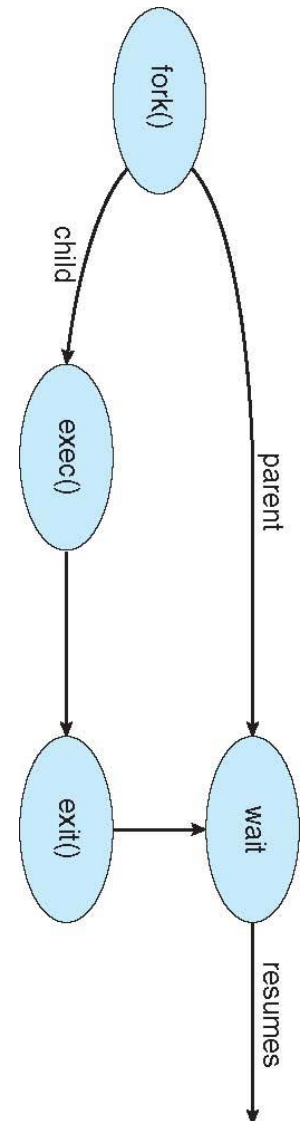
```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```





Operations on Processes



- Creating a Separate Process via Windows API

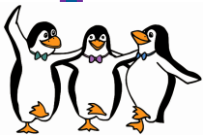
```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

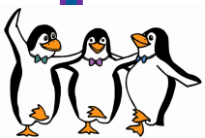




Operations on Processes

■ Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
 - ▶ Returns status data from child to parent (via `wait()`)
 - ▶ Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the `abort()` system call.
 - ▶ Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates





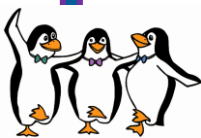
Operations on Processes

■ Process Termination

- Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - ▶ This is called **cascading termination**. All children, grandchildren, etc. are terminated.
 - ▶ The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the **pid** of the terminated process

```
pid = wait(&status);
```

- When a process terminates, its resources are deallocated by the OS.
- But, its entry in the process table must remain there until the parent calls `wait()`.
 - ▶ If the parent did not invoke `wait()`, the process is called a **zombie**
 - ▶ If parent terminated without invoking `wait()`, process is an **orphan**
 - The *init* process becomes the parent and issues `wait()` periodically to clear the orphans





Operations on Processes

■ Multiprocess Architecture – Chrome Browser

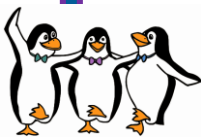
- Many web browsers ran as single process (some still do)
 - ▶ If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - ▶ **Browser** process manages user interface, disk and network I/O
 - ▶ **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - ▶ **Plug-in** process for each type of plug-in





Interprocess Communication

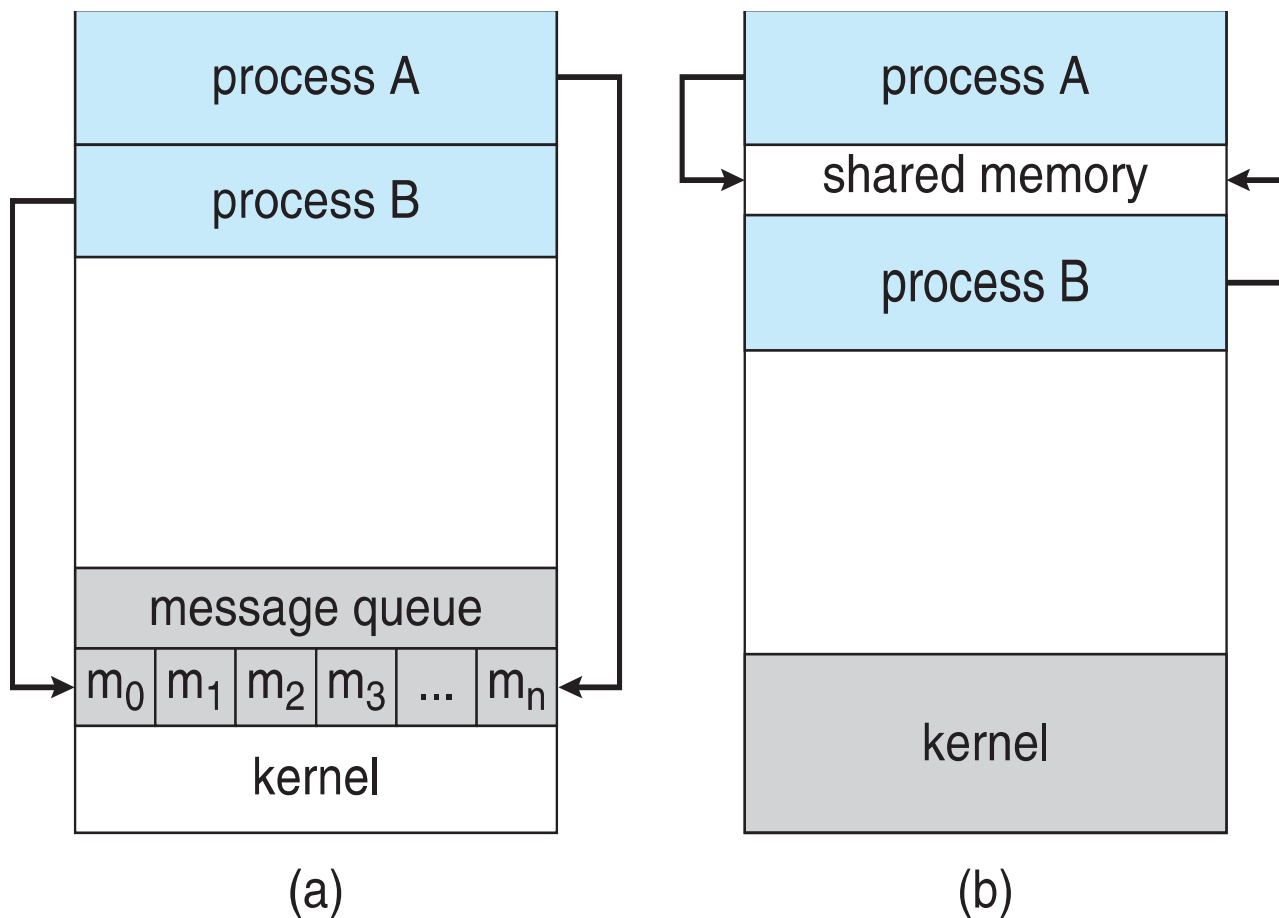
- Processes within a system may be **independent** or **cooperating**
 - **Independent process** cannot affect or be affected by the execution of another process
 - **Cooperating process** can affect or be affected by other processes, including sharing data
 - ▶ **Reasons for cooperating processes:**
 - **Information sharing** (e.g. shared file)
 - **Computation speedup**: break the process into multiple tasks, you need a multicore processor.
 - **Modularity**: divide the system functions into separate processes or threads
 - **Convenience**: a user may work on many tasks at the same time
 - ▶ Cooperating processes need **interprocess communication (IPC)** mechanism to allow them exchange data and information
 - ▶ Two models of IPC
 - **Shared memory**
 - **Message passing**



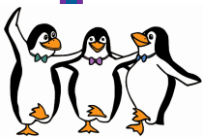


Interprocess Communication

- Communications Models



(a) Message passing. (b) shared memory.

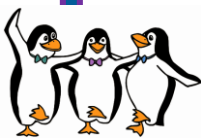




Interprocess Communication

■ Shared-Memory Systems

- An area of memory shared among the processes that wish to communicate
 - ▶ Typically, a shared-memory region resides in the address space of the process creating it.
 - Other processes must attach it to their address space
 - ▶ The communication is under the control of the users processes not the operating system.
 - ▶ Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.



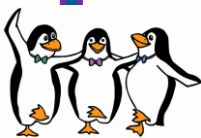


Interprocess Communication

■ Shared-Memory Systems

● **Producer-Consumer Problem**

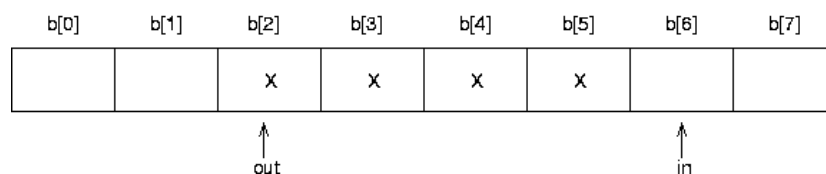
- ▶ It is a common paradigm for cooperating processes,
 - *producer* process produces information that is consumed by a *consumer* process
 - For example, a compiler may produce assembly code that is consumed by an assembler
 - To allow producer and consumer processes to run concurrently, we must have available a buffer of items that can be filled by the producer and emptied by the consumer.
 - **Two types of buffers can be used:**
 - » **unbounded-buffer** places no practical limit on the size of the buffer
 - » The consumer may have to wait for new items, but the producer can always produce new items
 - » **bounded-buffer** assumes that there is a fixed buffer size
 - » The consumer must wait if the buffer is empty, and the producer must wait if the buffer is full.





- The shared buffer is implemented as a **circular array** with two logical pointers: **in** and **out**.
 - ▶ The variable **in** points to the next free position in the buffer; **out** points to the first full position in the buffer.
 - ▶ The buffer is **empty** when **in == out**;
 - ▶ The buffer is **full** when **((in + 1) % BUFFER SIZE) == out**.
- The following variables reside in a region of memory shared by the producer and consumer processes

```
int out = 0;
```



-



Interprocess Communication

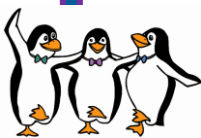
■ Bounded-Buffer – Shared-Memory Solution

● Producer code

```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

● Consumer code

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next consumed */
}
```

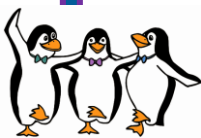




Interprocess Communication

■ Message-Passing Systems

- Mechanism provided by the OS for processes to communicate and to synchronize their actions without sharing the same address space
- It is useful in a distributed environment, where the communicating processes may reside on different computers connected by a network.
 - ▶ For example, an Internet chat program
- Message system – processes communicate with each other without resorting to shared variables
- Message-passing facility provides at least two operations:
 - ▶ **send**(*message*)
 - ▶ **receive**(*message*)
- The *message* size is either fixed or variable

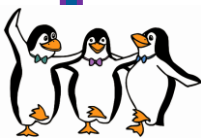




Interprocess Communication

■ Message-Passing Systems

- If processes P and Q wish to communicate, they need to:
 - ▶ Establish a **communication link** between them
 - ▶ Exchange messages via send/receive
- Implementation issues:
 - ▶ How are links established?
 - ▶ Can a link be associated with more than two processes?
 - ▶ How many links can there be between every pair of communicating processes?
 - ▶ What is the capacity of a link?
 - ▶ Is the size of a message that the link can accommodate fixed or variable?
 - ▶ Is a link unidirectional or bi-directional?





Interprocess Communication

■ Message-Passing Systems

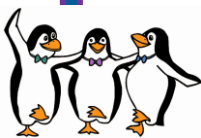
● Implementation of communication link

▶ Physical implementation:

- Shared memory
- Hardware bus
- Network

▶ Logical implementation:

- Direct or indirect communication
- Synchronous or asynchronous communication
- Automatic or explicit buffering





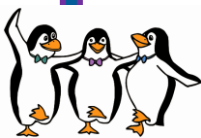
Interprocess Communication

■ Message-Passing Systems

- **Naming**: can be direct or indirect communication

- ▶ **Direct communication**

- Under direct communication, processes must name each other explicitly:
 - » **send** (P , *message*) – send a message to process P
 - » **receive**(Q , *message*) – receive a message from process Q
- **Properties of communication link**
 - » Links are established automatically: each process need to know each other identity
 - » A link is associated with exactly one pair of communicating processes
 - » Between each pair there exists exactly one link
 - » The link may be unidirectional, but is usually bi-directional



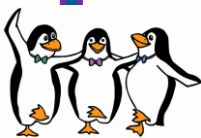


Interprocess Communication

■ Message-Passing Systems

● Indirect communication

- ▶ With indirect communication, messages are sent to and received from **mailboxes** (also referred to as **ports**)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- ▶ The send() and receive() primitives are defined as follows:
 - **send(A, message)**: Send a message to mailbox A.
 - **receive(A, message)**: Receive a message from mailbox A.
- ▶ **Properties of communication link**
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links, with each link corresponding to one mailbox
 - Link may be unidirectional or bi-directional

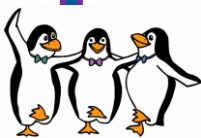




Interprocess Communication

■ Message-Passing Systems

- A mailbox that is owned by the operating system is independent and is not attached to any particular process.
- The operating system then must provide a mechanism that allows a process to do the following:
 - ▶ Create a new mailbox (port)
 - ▶ Send and receive messages through mailbox
 - ▶ Delete a mailbox
- The process that creates a new mailbox is its owner and is the only one that can receive messages through this mailbox.
 - ▶ However, the ownership and receiving privilege may be passed to other processes through appropriate system calls.

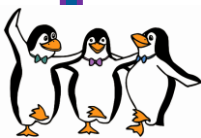




Interprocess Communication

■ Message-Passing Systems

- Mailbox sharing
 - ▶ Suppose that processes P_1 , P_2 , and P_3 share mailbox A
 - ▶ P_1 sends; P_2 and P_3 receive
 - ▶ Who gets the message?
- Solutions:
 - ▶ Allow a link to be associated with at most two processes
 - ▶ Allow only one process at a time to execute a receive operation
 - ▶ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.





Interprocess Communication

■ Synchronization

- There are different design options for implementing *send()* and *receive()* primitives:

- ▶ **Message passing may be either **blocking** or **non-blocking****

- **Blocking** is considered **synchronous**

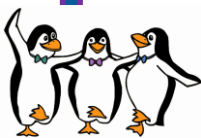
- » **Blocking send:** the sender is blocked until the message is received
- » **Blocking receive:** the receiver is blocked until a message is available

- **Non-blocking** is considered **asynchronous**

- » **Non-blocking send:** the sender sends the message and resumes operation
- » **Non-blocking receive:** the receiver receives either a valid message, or null message

■ Different combinations possible

- » If both send and receive are blocking, we have a **rendezvous** between the sender and the receiver.





Interprocess Communication

■ Synchronization

- Producer-consumer becomes trivial when using blocking send() and receive()

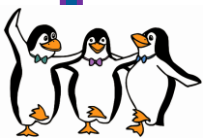
- The producer process using message passing.

```
message next_produced;
while (true) {
    /* produce an item in next produced */
    send(next_produced);
}
```

- The consumer process using message passing.

```
message next_consumed;
while (true) {
    receive(next_consumed);

    /* consume the item in next consumed */
}
```

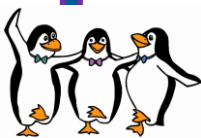




Interprocess Communication

■ Buffering

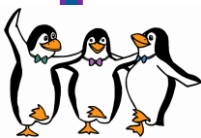
- Whether communication is direct or indirect, messages exchanged by communicating processes reside in a temporary queue.
- **Such queues can be implemented in three ways:**
 1. **Zero capacity:** no messages are queued on a link.
 - Sender must wait for receiver (rendezvous)
 2. **Bounded capacity:** queue has finite length of n messages
 - Sender must wait if link is full, otherwise can resume operation
 3. **Unbounded capacity:** infinite length
 - Sender never waits





Examples of IPC Systems - POSIX

- Several IPC mechanisms are available for POSIX systems,
 - shared memory and message passing
- POSIX shared memory is organized using memory-mapped files, which associate the region of shared memory with a file
- POSIX Shared Memory
 - Process first creates shared memory segment
`shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
 - Also used to open an existing segment to share it
 - Set the size of the object
`ftruncate(shm_fd, 4096);`
 - Now the process could write to the shared memory
`printf(shared_memory, "Writing to shared memory");`





Examples of IPC Systems - POSIX

- IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

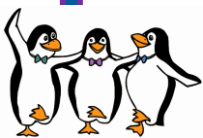
    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```





Examples of IPC Systems - POSIX

- IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

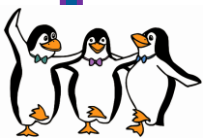
    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

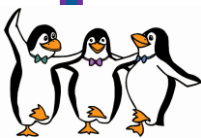
    return 0;
}
```





Examples of IPC Systems - Mach

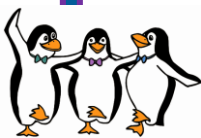
- Mach communication is message based
 - Even system calls are made by messages
 - Each task gets two mailboxes at creation- Kernel and Notify
 - Only three system calls needed for message transfer
`msg_send()` , `msg_receive()` , `msg_rpc()`
 - Mailboxes needed for communication, created via
`port_allocate()`
 - Send and receive are flexible, for example four options if mailbox full:
 - ▶ Wait indefinitely until there is room in the mailbox
 - ▶ Wait at most n milliseconds
 - ▶ Return immediately
 - ▶ Temporarily cache a message





Examples of IPC Systems - Windows

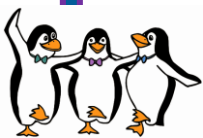
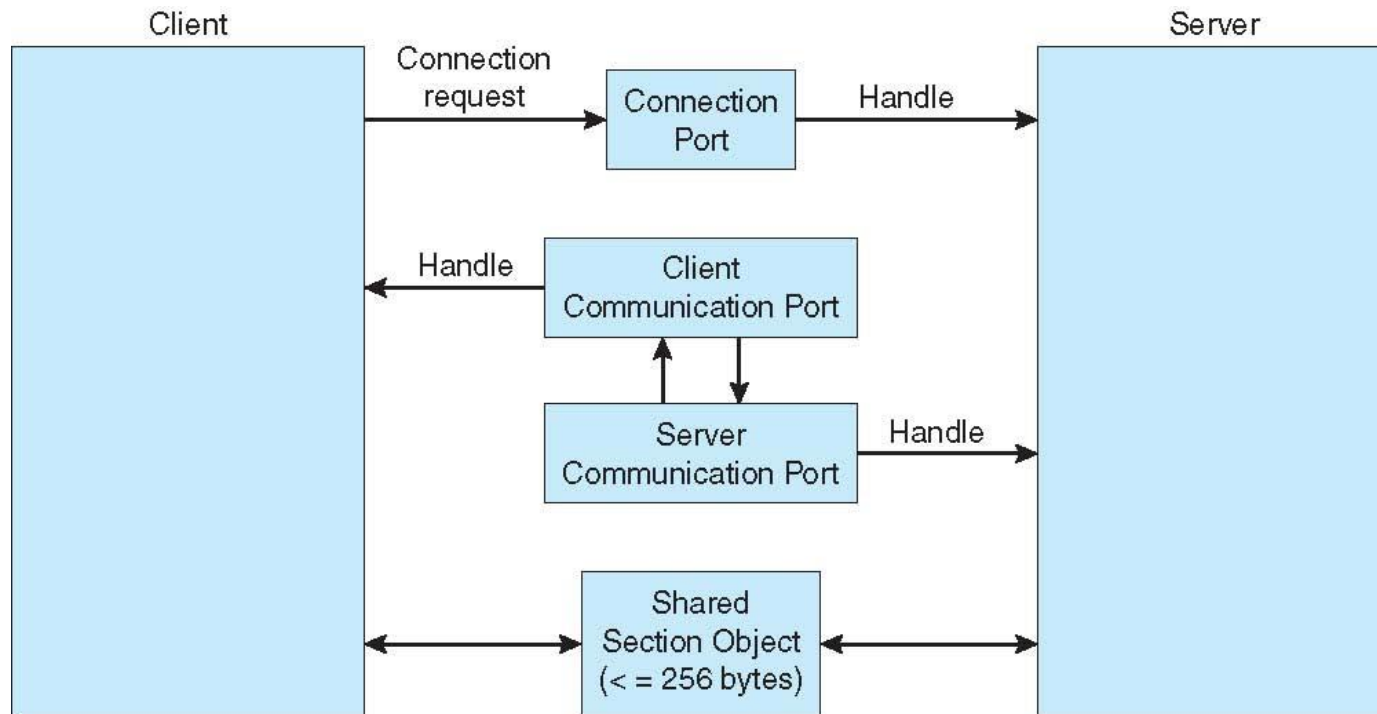
- Message-passing centric via **advanced local procedure call (LPC)** facility
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels
 - Windows uses two types of ports: **connection ports** and **communication ports**.
 - Communication works as follows:
 - ▶ The client opens a handle to the subsystem's **connection port** object.
 - ▶ The client sends a connection request.
 - ▶ The server creates two private **communication ports** and returns the handle to one of them to the client.
 - ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.





Examples of IPC Systems - Windows

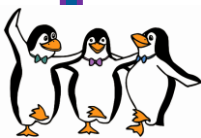
- Local Procedure Calls in Windows





Communications in Client-Server Systems

- The shared memory and message passing can be used for communication in client–server systems
- There are other strategies for communication in client–server systems
 - Sockets
 - Remote Procedure Calls
 - Pipes
 - Remote Method Invocation (Java)

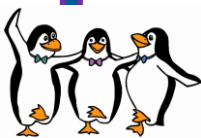




Communications in Client-Server Systems

■ Sockets

- A **socket** is defined as an endpoint for communication
- A pair of processes communicating over a network employs a pair of sockets
- It is identified by the concatenation of **IP address** and **port** – a number included at start of message packet to differentiate network services on a host
 - ▶ The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
 - ▶ All connections must be unique
- All ports below 1024 are **well known**, used for standard services
 - ▶ FTP server listens to port 21; a web, or HTTP, server listens to port 80
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running





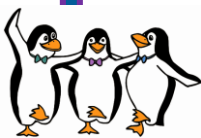
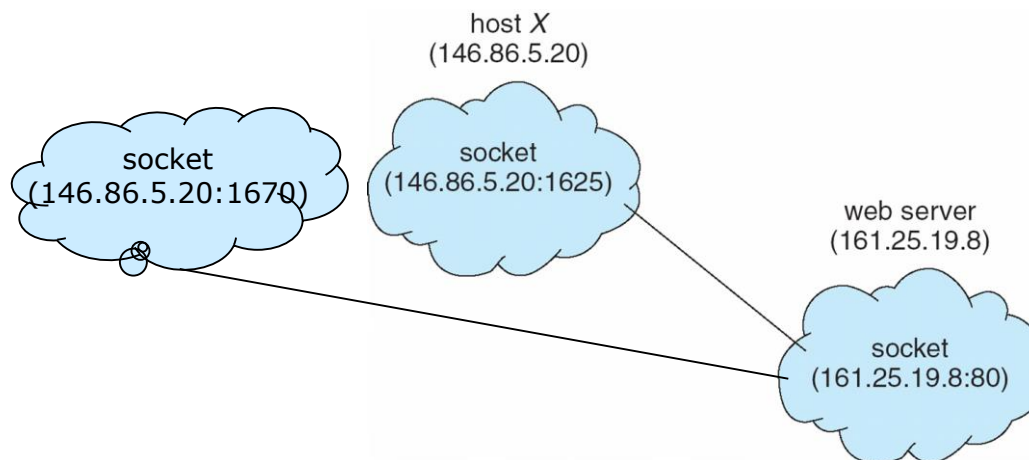
Communications in Client-Server Systems

■ Sockets

- Sockets use a client–server architecture: The server waits for incoming client requests by listening to a specified port
- When a client process initiates a request for a connection, it is assigned a port by its host computer.

▶ Example,

- If a client on host X with IP address 146.86.5.20 wishes to establish a connection with a web server (which is listening on port 80) at address 161.25.19.8, host X may be assigned port 1625.
- The connection consists of a pair of sockets: (146.86.5.20:1625) on host X and (161.25.19.8:80) on the web server.

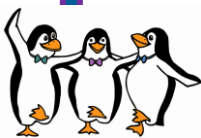




Communications in Client-Server Systems

■ Sockets in Java

- Three types of sockets
 - ▶ **Connection-oriented (TCP)**: implemented with the **Socket** class
 - ▶ **Connectionless (UDP)**: use the **DatagramSocket** class
 - **MulticastSocket** class: is a subclass of the **DatagramSocket** class.
 - » data can be sent to multiple recipients
 - ▶ **Example:**
 - Date server that uses connection-oriented TCP sockets.
 - » The operation allows clients to request the current date and time from the server.





Communications in Client-Server Systems

■ Sockets in Java

- Consider this “Date” server:
- The operation allows clients to request the current date and time from the server.
- The server creates a `ServerSocket` that specifies that it will listen to port **6013**. The server then begins listening to the port with the **`accept()`** method. The server blocks on the `accept()` method waiting for a client to request a connection. When a connection request is received, `accept()` returns a socket that the server can use to communicate with the client.

```
import java.net.*;
import java.io.*;

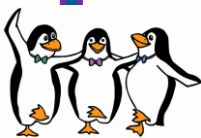
public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





Communications in Client-Server Systems

■ Sockets in Java

- Consider this “Date” client:
- A client communicates with the server by creating a socket and connecting to the port on which the server is listening.
- The client creates a Socket and requests a connection with the server at IP address **127.0.0.1** on port **6013**.
- Once the connection is made, the client can read from the socket using normal stream I/O statements.
- Close the connection

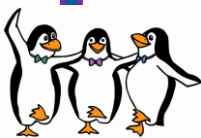
```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

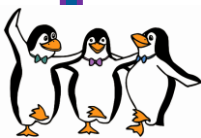




Communications in Client-Server Systems

■ Remote Procedure Calls (RPC)

- RPC abstracts procedure calls between processes on networked systems
 - ▶ It uses ports for service differentiation
- A **port** is a number included at the start of a message packet
- The RPC system hides the details that allow communication to take place by providing a **stub** on the client side.
 - ▶ When the client invokes a remote procedure, the RPC system calls the appropriate stub
 - ▶ The client-side stub locates the server and **marshalls** the parameters
 - ▶ The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**

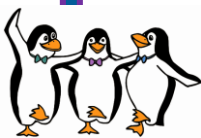




Communications in Client-Server Systems

■ Remote Procedure Calls (RPC)

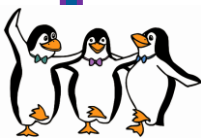
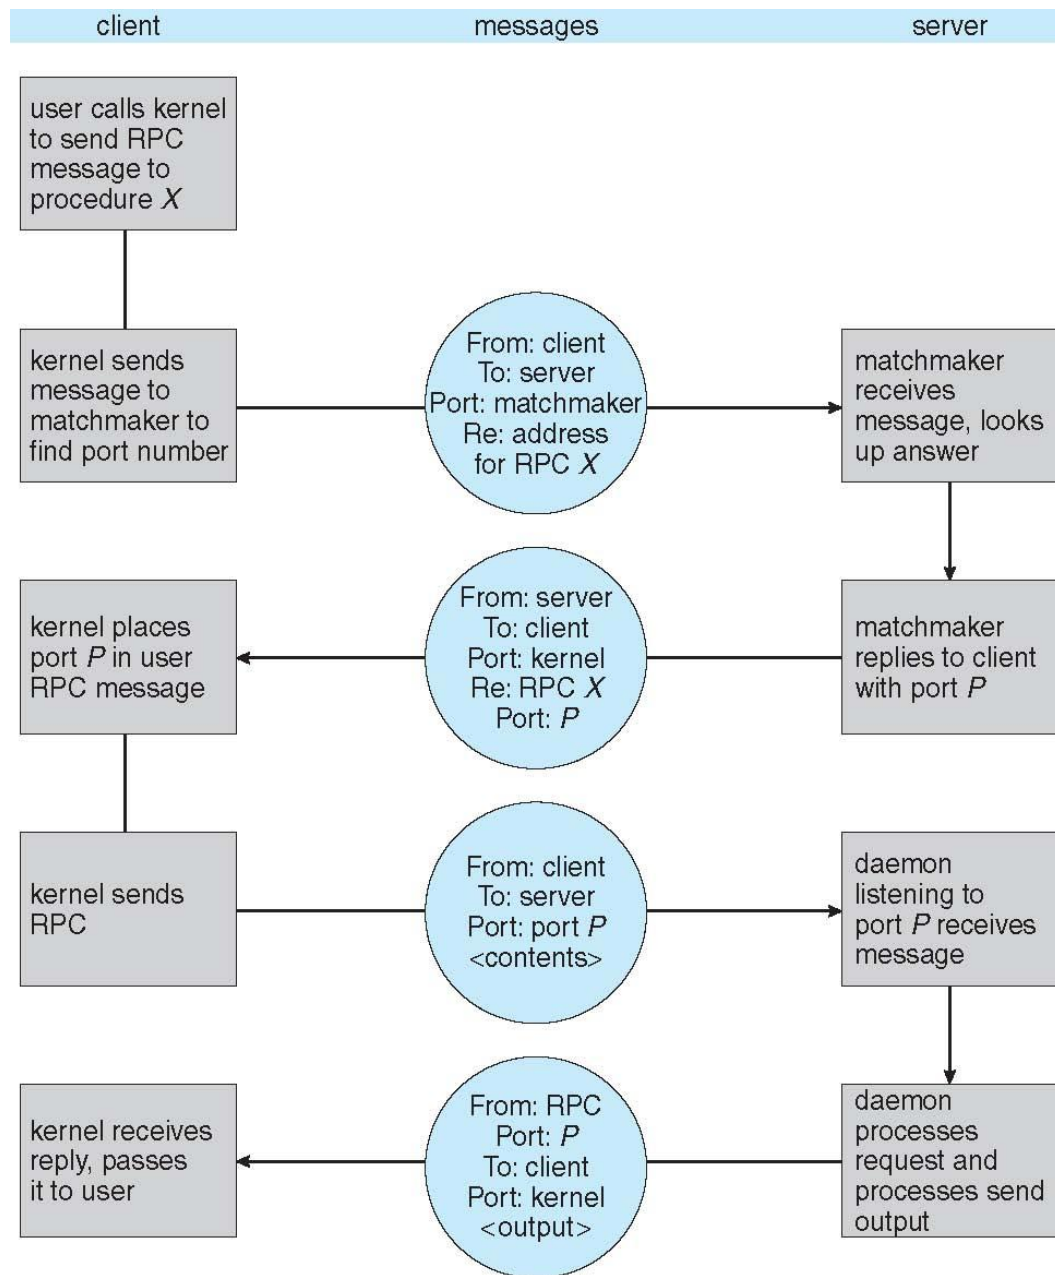
- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
 - ▶ **Big-endian** and **little-endian**
- Remote communication has more failure scenarios than local
 - ▶ Messages can be delivered **exactly once** rather than **at most once**
- OS typically provides a rendezvous (or **matchmaker**) service to connect client and server





Communications in Client-Server Systems

- Execution of RPC

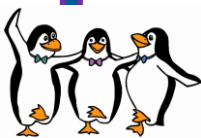




Communications in Client-Server Systems

■ Pipes

- Acts as a conduit allowing two processes to communicate
- In implementing a pipe, four issues must be considered::
 - ▶ Is communication unidirectional or bidirectional?
 - ▶ In the case of two-way communication, is it half or full-duplex?
 - ▶ Must there exist a relationship (such as, **parent-child**) between the communicating processes?
 - ▶ Can the pipes be used over a network?
- **Ordinary pipes:** cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- **Named pipes:** can be accessed without a parent-child relationship.

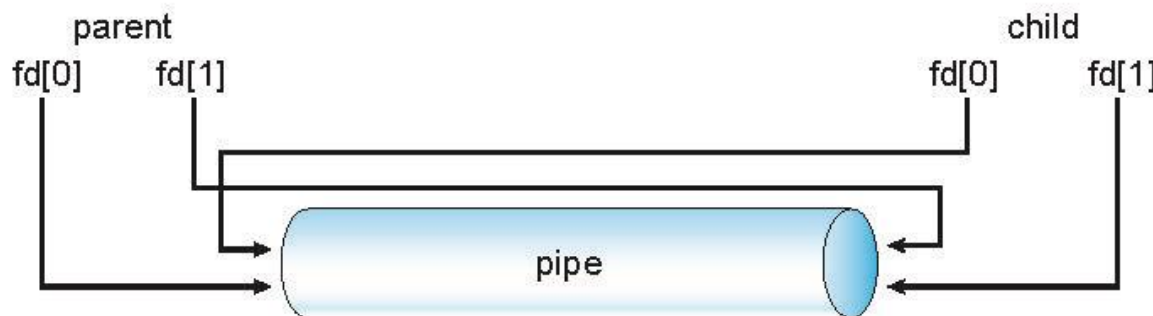




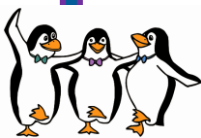
Communications in Client-Server Systems

■ Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore **unidirectional**
- Require parent-child relationship between communicating processes



- Windows calls these **anonymous pipes**





Communications in Client-Server Systems

■ Ordinary Pipes

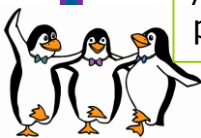
- On UNIX systems, ordinary pipes are constructed using the function

pipe(int fd[])

- It creates a pipe that is accessed through the int fd[] file descriptors: fd[0] is the read-end of the pipe, and fd[1] is the write-end.
- pipes can be accessed by read() and write() system calls just like files

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#define BUFFER SIZE 25
#define READ END 0
#define WRITE END 1
int main(void)
{
    char write msg[BUFFER SIZE] = "Greetings";
    char read msg[BUFFER SIZE];
    int fd[2];
    pid_t pid;
    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }
    /* fork a child process */
    pid = fork();
```

```
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    if (pid > 0) { /* parent process */
        /* close the unused end of the pipe */
        close(fd[READ END]);
        /* write to the pipe */
        write(fd[WRITE END], write msg, strlen(write msg)+1);
        /* close the write end of the pipe */
        close(fd[WRITE END]);
    }
    else { /* child process */
        /* close the unused end of the pipe */
        close(fd[WRITE END]);
        /* read from the pipe */
        read(fd[READ END], read msg, BUFFER SIZE);
        printf("read %s", read msg);
        /* close the write end of the pipe */
        close(fd[READ END]);
    }
    return 0; }
```

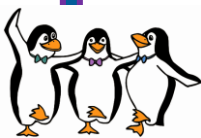




Communications in Client-Server Systems

■ Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems

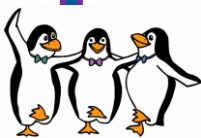




Communications in Client-Server Systems

■ Named Pipes

- On **UNIX**:
 - ▶ Referred to as FIFOs.
 - ▶ Once created, they appear as typical files in the file system.
 - ▶ A FIFO is created with the **mkfifo()** system call and manipulated with the ordinary **open()**, **read()**, **write()**, and **close()** system calls.
 - ▶ They are half-duplex. For full duplex use two FIFOs
 - ▶ Can be used in the same machine only, otherwise use sockets for inter machine communication
 - ▶ Example: **ls | more**
 - Setting up a pipe (**|**) between the **ls** and **more** commands (which are running as individual processes)
 - The **ls** command serves as the producer, and its output is consumed by the **more** command





Communications in Client-Server Systems

■ Named Pipes

- On **Windows**:

- ▶ Full-duplex communication is allowed, and the communicating processes may reside on either the same or different machines.
- ▶ created with the **CreateNamedPipe()** function, and a client can connect to a named pipe using **ConnectNamedPipe()**.
- ▶ Communication over the named pipe can be accomplished using the **ReadFile()** and **WriteFile()** functions.
- ▶ Example: **dir | more**

